


A summary of all the
good, the bad and the ugly



Outline

- 1 Outliers
 - Detecting outliers
 - Sensitivity to outliers: Discrete variables
 - Sensitivity to outliers: Continuous variables
- 2 Some wrapper functions
- 3 Summary of concepts
 - The need for a vocabulary of distributions
 - A simple simulation
- 4 Summary

Learning Objectives


- 1 Learn methods used to identify outliers.
- 2 Understand simple wrapper functions.
- 3 Appreciate the key concepts discussed in this set of e-learning materials.

Outliers

Detecting outliers

Detecting outliers: Examine empirical and theoretical quantiles

- Examine the *empirical* and *theoretical* quantiles of our data.
 - ▶ In selecting a distribution $p(x|\theta)$ or $f(x|\theta)$, we assume that the population data follows the distribution.
 - ▶ E.g., Q-Q plot for `seatbelts_4$PetrolPrice` fitted to the normal distribution.





Detecting outliers: The 1.5 IQR rule

- The 1.5 (interquartile range) IQR rule.
 - ▶ We can use the `boxplot()` function to make boxplots.

```
boxplot(seatbelts_4$VanKilled,  
        main = "Boxplot for VanKilled", ylab="Number of van drivers")  
boxplot(seatbelts_4$PetrolPrice,  
        main = "Boxplot for PetrolPrice", ylab="Petrol prices")
```





- ▶ Outliers are taken to be data points 1.5 IQR from the first and third quartiles.



Detecting outliers: The 1.5 IQR rule

cont'd



- Let us now remove the outliers.



Sensitivity to outliers: Discrete variables

Sensitivity to outliers: Discrete variables



- Let us revisit our example involving `seatbelts_4$VanKilled`.
- For each model, let us briefly compare a pair of rootograms.
- Let us start with the *Poisson* model.



Sensitivity to outliers: Discrete variables

cont'd

- Let us do the same for the *binomial* model.




Sensitivity to outliers: Discrete variables


cont'd

- Let us do the same for the *negative binomial* and *geometric* models.
- We will not always see an improvement upon removing outliers.
- In many cases, outliers can be important.


Negative binomial (with outliers)




Negative binomial (outliers removed)



Geometric (with outliers)



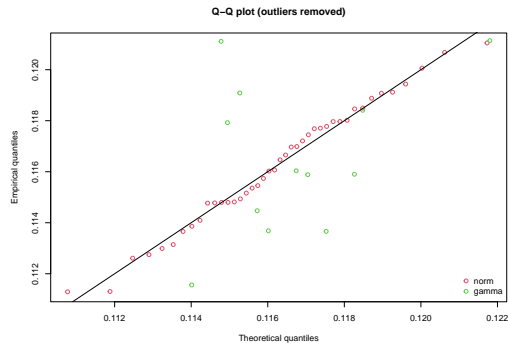
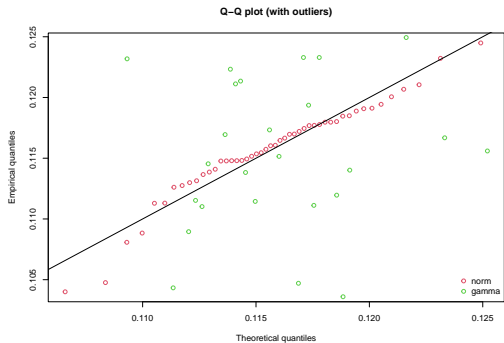
Geometric (outliers removed)



Sensitivity to outliers: Continuous variables

Sensitivity to outliers: Continuous variables

- Let us revisit our example involving `seatbelts_4$PetrolPrice`.
- For each model, let us plot a Q-Q plot.
- In the absence of outliers, model-fit for the normal distribution (in red) improved slightly.
- Outliers do not always impact the end result in an huge way.
- Outliers should be treated on a case-to-case basis.



Some wrapper functions

Some wrapper functions: a simple example

- In the following weeks, we will be using several wrapper functions.
- You are not expected to understand *all* components of these functions.
- Aim to be comfortable using these functions.
- Recall the following code:

```
norm_petrol <- fitdist(data = seatbelts_4$PetrolPrice ,  
                      distr = "norm")  
qqcomp(norm_petrol)
```

- What if you want to choose a different model?
- How about a different dataset and variable?
 - ▶ You have to change the `distr` argument.
- What if you have many more steps?

Some wrapper functions: a simple example

cont'd

- We can create a custom function that is like a template.
- The following function will carry out two steps sequentially.

```
qqplot_dist <- function(vec, distname){  
  fitted <- fitdist(data = vec, # Set up the data  
                    distr = distname) # Select the model  
  qqcomp(fitted) # Plot Q-Q plot  
}
```

- This function is very similar to the earlier syntax:

```
norm_petrol <- fitdist(data = seatbelts_4$PetrolPrice,  
                      distr = "norm")  
qqcomp(norm_petrol)
```

- To obtain the Q-Q plot from before, we have to correctly specify the arguments `vec` and `distname`.

```
qqplot_dist(vec = seatbelts_4$PetrolPrice,  
            distname = "norm")
```



Some wrapper functions: a simple example (cont'd)

- We can then easily plot the Q-Q plot for a different variable, say, for `kms`.

```
qqplot_dist(vec = seatbelts_4$kms ,  
            distname = "norm")
```

- We can do the same for a different model, say, for the *Cauchy* distribution.

```
qqplot_dist(vec = seatbelts_4$kms ,  
            distname = "cauchy")
```



Some wrapper functions: a slightly more complicated example

- We can modify the `qqplot_dist()` function so that we can plot a Q-Q plot for several models at once.


```
qqplot_dist_multi <- function(vec, distname_multi){  
  vec <- as.vector(vec) # Ensure that vec is a vector  
  distname_multi <- as.list(as.character(distname_multi))  
  fitted <- lapply(distname_multi, fitdist, data=vec)  
  qqcomp(fitted)  
}
```

Some wrapper functions: a slightly more complicated example

cont'd

- To plot the Q-Q plot for both exponential and normal distributions, we can use this function.

```
qqplot_dist_multi(vec = seatbelts_4$PetrolPrice ,  
                  distname_multi = list("exp", "norm"))
```



Summary of concepts

The need for a vocabulary of distributions

The need for a vocabulary of distributions

The good:

- We have seen several distributions.
- Strive to familiarise yourself with more common distributions.

The bad:

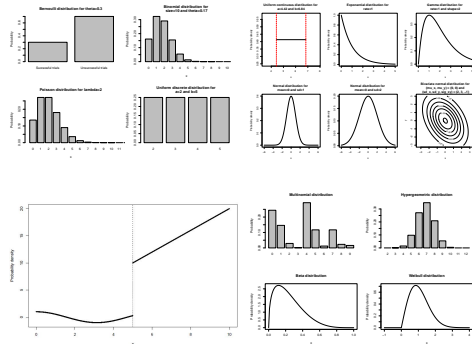
- Not all distributions can be mathematically formalised.
- Not always possible to describe a distribution in terms of parameters and the support.

The ugly:

- When encountering an unfamiliar distribution, focus on its parameters and support.

Looking forward:


- We will be using different distributions as building blocks for various simulation models.



A simple simulation

A simple simulation

- Suppose there are 4 students under a certain special programme.
 - ▶ Past experience tells us that 40% of students will get the bursary without extra help.
 - ▶ Let us then take 0.4 as the *empirical* probability of getting the bursary.
 - ▶ Each student can be treated as an individual Bernoulli trial.
 - ▶ Then, we expect 1.6 out of 4 students to obtain the bursary.
- Let us further suppose that we want to increase every student's chance of getting the bursary.
 - ▶ There are exactly two students who are helpable.
 - ★ If one-to-one tuition is provided, the probability of obtaining the bursary for these students can be increased to 0.8.
 - ▶ However, there are only two tutors available.
- Is the strategy of providing tuition to *two randomly selected students* a good strategy?



A simple simulation: Using a wrapper function

- We shall use a wrapper function called `sim_students()` to simulate selecting two students randomly.

```
sim_students <- function(outcomes=FALSE){  
  students <- c("a", "b", "c", "d")  
  selected.students <- sample(x = students, size = 2)  
  pa <- 0.4;pb <- 0.4;pc <- 0.4;pd <- 0.4  
  if(selected.students[1] == "a"){pa <- 0.8}  
  if(selected.students[2] == "b"){pb <- 0.8}  
  if(selected.students[1] == "b"){pb <- 0.8}  
  if(selected.students[2] == "a"){pa <- 0.8}  
  if(outcomes == FALSE){  
    sum(rbinom(4, size = 1, prob = c(pa, pb, pc, pd)))  
  }else{rbinom(4, size = 1, prob = c(pa, pb, pc, pd))}  
}
```

A simple simulation: Running 1 simulation

- Let us run the simulation once, and set the seed as 1.
 - ▶ Setting `outcomes = TRUE` allows us to view the outcomes.

```
set.seed(1)
sim_students(outcomes = TRUE)
```

```
## [1] 0 0 1 1
```

- ▶ Let us run this again, this time without setting up the argument.

```
set.seed(1)
sim_students()
```

```
## [1] 2
```

A simple simulation: Running 100 simulations

- Using the `replicate()` function, let us now run the simulation 100 times, and set the seed as 123.

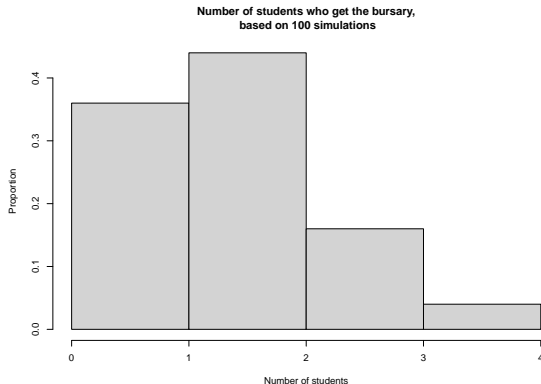
```
set.seed(123)
sim_100 <- replicate(n = 100,
                      sim_students())
```

- Next, we can compute the mean number of students who got the bursary.

```
mean(sim_100)
```

```
## [1] 1.84
```

- This means that 1.84 out of 4 students are expected to obtain the bursary with extra help given.
- Let us also plot the histogram to get a better idea on how the simulated values are distributed.



A simple simulation: Running 100 000 simulations

- Law of large numbers: running more simulations will bring the empirical probabilities closer to their theoretical counterparts.

- Let us now run 100 000 simulations.

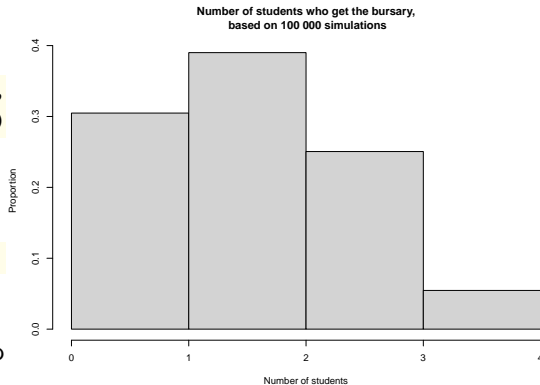
```
sim_100000 <- replicate(n = 100000,  
                        sim_students())
```

- Again, we can compute the mean number of students who got the bursary.

```
mean(sim_100000)
```

```
## [1] 2.00241
```

- This means that 2 out of 4 students are expected to obtain the bursary with extra help given.
- Let us also plot the histogram to get a better idea on how the simulated values are distributed.



Summary

Summary

In this video, we have:

- Learned some methods used to identify outliers.
- Learned to use simple wrapper functions.
- Discussed the good, bad and ugly aspects of model-fitting.
- Discussed a simple simulation.

References



R-data — seatbelts dataset.



Hoaglin, D. C., Iglewicz, B., and Tukey, J. W. (1986).
Performance of some resistant rules for outlier labeling.
Journal of the American Statistical Association, 81(396):991–999.



Wasserman, L. (2004).
All of statistics springer new york.