


Introduction to Permutation Tests



Outline

- 1 Permutation Tests
- 2 Implement A Permutation Test in R
- 3 Summary

Permutation Tests

A new approach to statistical inference:

- ***Permutation-based hypothesis testing***, which is also called "hypothesis testing with randomisation".
- We will simply refer to the tests as "***permutation tests***".

Learning Objectives

In this video, we will

- Elaborate the idea of permutation tests.
- Implement a permutation test in *R*.

Introduction to Permutation Tests

Why are we interested in permutation tests?

- 1 Permutation tests allow us to use any test statistic for which the sampling distribution has not been derived yet.
 - ▶ In the two-sample set-up, examples include the difference of medians, ratio of means, and ratio of variances.
- 2 Permutation tests require less distributional assumptions.
 - ▶ We do not need to assume a particular family of distributions for the population data.
 - ▶ In particular, permutation tests work well in the cases when the sample size is small, or the population data don't follow a normal distribution.

How to implement permutation tests?

- Randomised experiment: 20 students were randomly split into two groups: A and B.
- H_0 : The population mean scores of the two groups are the same.
- If H_0 is true, the difference between the two samples' mean scores should be due to natural variability inherent to the sample dataset.
- We will set up a sham randomised experiment by reallocating the group order.
- We will refer to the sham experiment's result as the "***permutation sample***".

Original Sample

Group	Scores	Group	Scores
A	64	B	60
A	66	B	69
A	50	B	84
...
A	59	B	61
A	65	B	66
A	43	B	54

“Shuffle the Cards”

Original Sample

Group	Scores	Group	Scores
A	64	B	60
A	66	B	69
A	50	B	84
...
A	59	B	61
A	65	B	66
A	43	B	54

Permutation Sample

Group	Scores	Group	Scores
A	64	A	60
B	66	B	69
A	50	B	84
...
B	59	B	61
A	65	A	66
A	43	B	54

“Shuffle the Cards”

Original Sample

Group	Scores	Group	Scores
A	64	B	60
A	66	B	69
A	50	B	84
...
A	59	B	61
A	65	B	66
A	43	B	54

Sorted Permutation Sample

Group	Scores	Group	Scores
A	64	B	66
A	60	B	69
A	50	B	84
...
A	66	B	61
A	65	B	59
A	43	B	54

Implement A Permutation Test in R

Some Preparations in R

- We first prepare the variables.

```
treatment <- df1$treatment  
outcome <- df1$outcome
```

```
> treatment  
[1] "a" "a" "a" "a" "a" "a" "a" "a" "a" "a" "b" "b" "b" "b" "b" "b" "b" "b"  
[19] "b" "b"  
> outcome  
[1] 64 66 50 36 43 60 66 59 65 43 60 69 84 81 48 59 64 61 66 54
```

- We can access one variable of the data frame, by specifying the name of the data frame, followed by a dollar sign and the variable name.

“sample()” Function

- The base R function “**sample()**” can help us reallocate the order of a variable or a vector.

```
set.seed(123)  
sample(1:5, size = 1)
```

```
[1] 3
```

```
sample(1:5, size = 5, replace=FALSE)
```

```
[1] 3 2 4 5 1
```

- It will randomly draw 5 numbers from 1 to 5, without replacement.
- This is equivalent to returning a permutation sequence of 1 to 5.

Generating a Permutation Sample

- We use the following code to generate a new group order for the permutation sample.

```
treatment_p <- sample(treatment, size= length(treatment), replace=
  FALSE)
treatment_p
```

```
[1] "a" "b" "b" "a" "b" "a" "a" "a" "a" "b" "b" "b" "a" "a" "a" "b" "b" "b"
[19] "a" "b"
```

- We define the test statistic as the difference of the two samples' means.
- We use the following code to calculate the test statistic of the permutation sample.

```
mean(outcome[treatment_p == "b"]) - mean(outcome[treatment_p == "a"])
```

```
[1] -6
```

```
outcome[treatment_p == "a"]
```

```
[1] 64 36 60 66 59 65 84 81 48 66
```

```
outcome[treatment_p == "b"]
```

```
[1] 66 50 43 43 60 69 59 64 61 54
```

“for loop”

- The syntax of **for loop** is as follows,

```
for (variable in sequence)
{
  statement
}
```

- Let us use the following example to explain how it works.

```
sum <- 0
for (i in 1:2) {
  sum <- sum + i
}
sum
```

```
[1] 3
```

- The example is iteratively calculating $1 + 2$.

Use “for loop” to Generate Multiple Permutation Samples

- Some preparations:

```
size <- 10000  
perm_dist <- rep(0, size)
```

- To generate 10,000 permutation samples:


```
for(i in 1:size){  
  # Take a permutation sample without replacement  
  treatment_p <- sample(treatment, size= length(treatment), replace  
    =FALSE)  
  # For each permutation sample, calculate the test statistic  
  perm_dist[i] <- mean(outcome[treatment_p == "b"]) -  
    mean(outcome[treatment_p == "a"])  
}
```

- The test statistic of each permutation sample is recorded at the i -th position of “perm dist”.

Generate the Histogram of all Test Statistics

- We use the “hist()” function to generate the histogram of the test statistics for the 10,000 permutation samples.


```
hist(perm_dist)
```



Generate the Histogram of all Test Statistics

Cont'd


- Next, we mark the original sample statistic, 9.4.
- ***p-value*** is defined as the proportion of the permutation test statistics that are at least as extreme as the observed 9.4.



Generate the Histogram of all Test Statistics

Cont'd

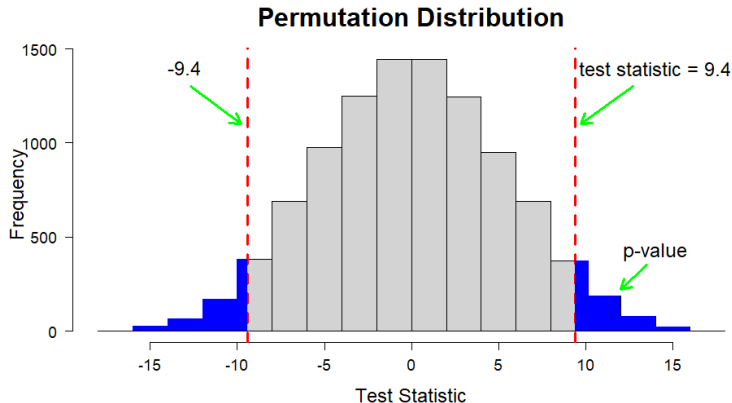
- Next, we mark the original sample statistic, 9.4.
- ***p-value*** is defined as the proportion of the permutation test statistics that are at least as extreme as the observed 9.4.



Generate the Histogram of all Test Statistics

Cont'd

- Next, we mark the original sample statistic, 9.4.
- ***p-value*** is defined as the proportion of the permutation test statistics that are at least as extreme as the observed 9.4.



Calculate p-value

- We can calculate the p-value as follows.

```
pvalue <- mean(abs(perm_dist) >= abs(original))  
pvalue
```

```
[1] 0.0762
```

```
head(abs(perm_dist) >= abs(original))
```

```
[1] FALSE TRUE TRUE FALSE TRUE FALSE
```

- Roughly 7.6% of the simulated permutation samples have more extreme values than the observed difference.
- As $0.076 > 0.05$, we do not reject the null hypothesis.

Define a Function

- The syntax of defining a function is as follows:

```
function name <- function(variable1, variable2)
{
  Actions
}
```

- We wrap the actions of one round of simulation inside the following function.

```
permutation.test <- function(treatment, outcome){
  # Generate a permutation sample
  treatment_p <- sample(treatment, size= length(treatment), replace
    =FALSE)
  # Calculate the test statistic for the permutation sample
  mean(outcome[treatment_p == "b"]) -
    mean(outcome[treatment_p == "a"])
}
```

- The function does two things: generate a permutation sample and calculate the test statistic.

Use the “replicate()” Function to Run Multiple Simulations

- To run the function once is equivalent to one simulation.
- Next, we use the “replicate()” function to run the function multiple times.

```
test <- replicate(10000, permutation.test(treatment, outcome))
```

- The Replicate function needs two inputs:
 - ▶ The front number indicates the number of simulations to be run.
 - ▶ The latter part quotes the function that we designed for one simulation.
- The following codes generate the permutation distribution and the p-value, just like the earlier ones.

```
hist(test)  
mean(abs(test) >= abs(original))
```

Check Assumption

Assumption of permutation tests

- 1 **Exchangeability:** The observations are exchangeable.
 - It means the labels of the dataset can be reordered, which does not affect the underlying joint distribution.
 - The assumption is naturally satisfied in the following two cases:
 - (a) a randomised experiment;
 - (b) a random sample without replacement.
 - As our example belongs to the former case, the assumption is verified.
 - In general, the exchangeability assumption is much weaker than those of the classical hypothesis tests.

Compare the p-values

- We tried the following code in the last video.

```
t.test(outcome~treatment, alternative = "two.sided", paired=F,  
var.equal=T, data = df1)
```

Two Sample t-test

data: outcome by treatment

t = -1.8767, df = 18, p-value = 0.07687

alternative hypothesis: true difference in means between group a and group b
is not equal to 0

95 percent confidence interval:

-19.923268 1.123268

sample estimates:

mean in group a	mean in group b
55.2	64.6

- In our case, the p-value of the permutation two sample test is very close to that of the classical two sample t-test.

Summary

Summary


We have:

- ▶ Explained the idea and the steps of conducting a permutation test.
- ▶ Learnt how to implement a permutation test in R.

In the next video,

We will use another case study to show how we deal with the more complex test statistic, and implement it in R.

References

-  Mine Çetinkaya-Rundel and Johanna Hardin (2021)
Introduction to Modern Statistics